

**MASARYK
UNIVERSITY**

FACULTY OF INFORMATICS

**Design and Implementation of a
Mobile Application with Offline
Support**

Master's Thesis

JIŘÍ LOUN

Brno, Fall 2025

MASARYK
UNIVERSITY

FACULTY OF INFORMATICS

**Design and Implementation of a
Mobile Application with Offline
Support**

Master's Thesis

JIRÍ LOUN

Advisor: RNDr. Pavel Novák

Department of Computer Systems and Communications

Brno, Fall 2025



Declaration

Hereby I declare that this paper is my original authorial work, which I have worked out on my own. All sources, references, and literature used or excerpted during elaboration of this work are properly cited and listed in complete reference to the due source.

Jiří Loun

Advisor: RNDr. Pavel Novák

Acknowledgements

TBA

Abstract

This thesis documents the process of designing and implementing an enterprise mobile application with offline support in a real-world scenario with specific functional requirements to be fulfilled. By analyzing viable solutions and architectures, this thesis also serves as a reference for architects to be consulted when designing an optimal approach for a similar application.

Keywords

Mobile, React, React Native, Kotlin, Offline mode, Architecture, Synchronization

Contents

Introduction	1
1 Offline-enabled mobile applications	2
1.1 Concept	2
1.2 Issues to be tackled	3
1.2.1 Offline data availability	3
1.2.2 Partial changes management	4
1.2.3 Local data management	6
1.2.4 Server synchronization	7
1.2.5 Offline authentication, authorization	9
2 Technical approaches	10
2.1 Frameworks	10
2.1.1 Native approaches	11
2.1.2 Cross-platform approaches	15
2.1.3 React PWA	20
2.1.4 Conclusion	23
2.2 Local data management	24
2.2.1 Request intercepting and processing	24
2.2.2 Local database	26
2.2.3 Hybrid solutions	27
2.3 Synchronization	28
2.3.1 Mitigation	28
2.3.2 Duplication	29
2.3.3 Attribute timestamping	29
2.3.4 CRDT	29
2.3.5 Advanced synchronization management	29
3 Implementation	30
3.1 Functional and non-functional requirements	30
3.2 Architecture, technology stack	30
3.3 Service Worker	30
3.4 Axios Interceptor	30
3.5 Local data management	30
3.6 Synchronization	30

4 Testing, issues, and future work	31
Conclusion	32
Bibliography	33
A Electronic attachments	35

List of Tables

List of Figures

Introduction

The primary objective of this thesis is to design and implement an offline-ready enterprise mobile client application. The secondary objective is to provide a comprehensive guide through designing such an application while considering different requirements. This guide features a list of available technological solutions with their advantages and setbacks while addressing the specific issues that arise in this architecture. The guide will be evaluated against a set of real-world functional and non-functional requirements and a real application that adheres to the result of that evaluation will be designed and implemented to be used by a customer as an enterprise solution.

The thesis consists of 4 chapters; the first chapter investigates the concept of offline-enabled mobile applications, their advantages, disadvantages, and specific challenges and architectural demands. The second chapter explores the approaches to be taken when designing an architecture and technology stack of such app - each approach is further documented, stating its strengths and weaknesses and practical scenarios in which each of them play their prime. Here the thesis not only focuses on different frameworks or languages, but also, the challenges introduced in the first chapter are confronted with different technical solutions and evaluated.

The third and fourth chapters focus on selecting the right approaches in different areas based on the functional and non-functional requirements, implementation of a mobile frontend client, testing, and other aspects of the practical execution of the project in the context of a real solution for a real customer.

The end application conforms with the requirements and is deployed to production for the end users to work with.

Describe motivation and purpose, business context ... TODO Tibi + nejaky clanek idealne? Why does the app solve the problem? What value does it bring?

1 Offline-enabled mobile applications

First, it is needed to take a look at offline-enabled enterprise mobile applications as a whole and identify the challenges they have to solve in addition to conventional mobile applications. By offline-enabled is meant an application that provides a significant part of its functionality without access to a server, while the extent of said functionality may vary based on a particular project and customer's needs.

What also tends to vary, and has a great influence on the project complexity, is the “extent” of offline capabilities. There may be applications that only need to display some data offline without modifying anything, or applications that need to survive only a couple of hours without synchronizing data with the server, but there may be applications that need to be backend-independent for days and then be able to synchronize the data with the server while running into as few conflicts as possible.

1.1 Concept

Mobile applications in the context of enterprise solutions always conform with the server-client schema — this thesis does not consider local-only applications as those have little to no use on the relevant market. In this server-client schema, there are two basic ideologies to pursue — server-heavy or client-heavy solution.

A server-heavy solution will rely largely on the backend capabilities. Any list filtering or sorting will be done through an endpoint, for example using a data search engine such as Elasticsearch¹. The app will store as little state as possible on the client and will depend on server-managed state, settings, or data.

If this concept is deepened further in the context of web applications, whole pages can be rendered on the server and sent already partially prepared (pre-rendered, waiting to be hydrated — enriched by client-rendered functionality — on the client) to the client's browser.

1. <https://www.elastic.co/elasticsearch>

Such approach is called SSR (Server Side Rendering) and is currently on the rise with frameworks such as Next.js² — a framework which has gained a lot of popularity in the recent years, according to Francisco Moretti[1].

A client-heavy application, on the other hand, will rather utilize the client's resources to process or transform data instead of asking the server to do it. It will prefer sorting or filtering data in the browser or managing state locally — for example, by using a state-management library like Redux³ or Zustand⁴.

When considering any offline capabilities, the application must not be heavily dependent on the server and thus will go in the client-heavy direction. To accommodate the offline needs, the concept of the application will need to further deepen the client's independence in order to be able to retain its functionality even when the server is not responsive due to signal loss. The specific problems that one encounters while developing such an app are further explored in the next section.

1.2 Issues to be tackled

The first and the most obvious task is to make the application available offline as a whole. While it may seem trivial for certain platforms — e.g. installing an Android application from Google Play Store — it may be the first obstacle in other approaches, e.g. React PWA (Progressive Web Application). These specifics will be examined in the following chapter; at present, we turn to more conceptual issues.

1.2.1 Offline data availability

In order for the application to be able to work with data in offline mode, these data must be downloaded in advance and kept updated. This presents a unique challenge that is not dependent on platform or

-
- 2. <https://nextjs.org/>
 - 3. <https://redux.js.org/>
 - 4. <https://github.com/pmndrs/zustand>

framework choice — it is, generally, not a good practice to download all data that could possibly be accessed in the application — that would practically mean data mirroring of whole database or a large part of it. That might not even be possible due to the database size. It is, therefore, quite imperative to try to limit the amount of data that should be available offline from the business logic side.

As an example, let's imagine a remote-synchronized enterprise calendar application for planning meetings. Meetings can be input from a device and synchronized to a server. All devices that are subscribed to a certain calendar (e.g. meetings for project X) have it automatically synced with the server to display events submitted from other devices. It might very well suffice to automatically download only the current week or month for offline purposes as there is probably low business value in viewing meeting plans several months in advance or in history.

1.2.2 Partial changes management

A related topic to the previous section is partial changes management — a concept of recording changes of data, not entire snapshots.

The use-case here is updating the local data to match the server state — the longer the application has been in offline mode, the higher the probability of the data being out of date. Specifically, the probability of the data being out of date after a certain period of time follows the Poisson's distribution[2]:

$$P(X = k) = \frac{\lambda^k}{k!} e^{-\lambda} \quad (1.1)$$

where λ is expected rate of occurrence and k is the number of occurrences. Let's assume that on average an instance of an entity is changed once every 2 hours during working hours by some user. That is 4 changes during a whole working day. Therefore $\lambda = 4$. The probability that the instance is still up to date the next day (after 8 working hours), in other words, $k = 0$, is approximately 1.83%.

Therefore, the local data should be periodically or manually updated to minimize the occurrence of synchronization conflicts. The

reason why the system should support partial changes management is that if there is none, the only option of a local data update is to bulk-download all data, as if there were none stored, and overwrite them.

There are two layers of this principle that a system can implement:

Tracking instances

The first layer is keeping track of whole objects (instances) that have been changed since a particular point in time — essentially, having an updated attribute in each instance that holds the timestamp of the last change of its attributes. In this case, when the client is asking the server for fresh data to update its local copy, it can only request fresh data of instances that were updated since the last data download.

There also has to be a check for any new relevant instances or any deleted ones - those must be deleted (or soft-deleted⁵) from the local data on the client.

Going back to the meeting calendar application example, the client would like to update its local list of planned meetings with their details. There would have to be an endpoint, say `POST /meetings/data/update` and the structure of the body of such a request could, for example, look like this:

```
{  
    "since": ISO string,  
    "knownIds": uuid [],  
}
```

where `since` is the timestamp of last data update on the device (every newer change needs to be provided) and `knownIds` is an array of IDs of planned meetings that the device currently tracks.

The response to such request could have this structure:

```
{  
    "added": Meeting [],
```

5. Soft-delete is a technique when an instance is marked as deleted instead of being deleted.

```
    "deleted": uuid[] ,  
    "updated": Meeting[] ,  
}
```

where in the updated and deleted parts, only the instances from the sent knownIDs are considered by the server.

Tracking attributes

Going even deeper in granularity, each attribute can have its timestamp and thus the response to the local-data-updating request may only contain changes in attributes and does not have to send the whole data object. While timestamping every attribute might sound like an overkill, as will be demonstrated later, in a complex offline-ready system, it may serve multiple purposes and therefore might be worth implementing.

Extending the simple request-response schema from the previous section, the only change would be that updated would be of type `MeetingPartial` - containing only meeting ID and attributes that have been changed.

1.2.3 Local data management

In a traditional client-server based application, the client should not persist any data that might be volatile. The aim is to fetch a fresh copy of data each time the user displays a certain page for them to work with the latest version of data⁶. That is especially important for data updates — in case there is any form of version control (even a simple one — every instance has a current version v and every PUT request must have version $v + 1$, otherwise the server refuses to process the request due to version inconsistency), the data freshness might be the differentiator between having the request processed or not.

After working with the data and possibly mutating them, they are sent to the server to process them and after receiving a positive

6. Here the realm of real-time applications is omitted as those are rather specific.

response, the data is refreshed. Hence, the local copy of data is fresh once again and the old copy is thrown away - as an example, let's imagine a updated attribute — one that stores the date and time of the last change that has been done to the instance. Updating the value of this attribute will only be done on the server, but the client definitely wants to see this attribute updated after submitting a mutation to the entity.

In the offline world, there is no fresh data and there is no server. All mutations have to be done on the local data copy, including attributes like updated. The server interaction needs to be faked by the application itself — which, in reality, means implementing a lot of server functionality on the client, which represents the client thickness.

1.2.4 Server synchronization

Synchronizing the changes made on the client to the server and managing potential conflicts is arguably one of the most difficult tasks to tackle in the field of offline-enabled applications. The task is to synchronize all changes made in offline mode to the server when coming back online as seamlessly as possible.

The reason why conflicts arise is that the system, and therefore (generally) any data in it can be accessed and mutated by other users, while one or more users are making changes in offline mode. There may be two general approaches:

State-based synchronization

All the mutating changes done in offline mode affect directly the local copy of data while marking it *dirty* — when switching back to online, every dirty instance is synchronized to the server in one mutation - one PUT request. Newly created (POST) instances that have later been updated are synchronized as one POST with all the updates already put in. Updating an instance and then deleting it means only the DELETE request is synchronized and the updates are thrown away.

Operation-based synchronization

Each and every update is propagated to the server. That means, there is a layer that tracks all the requests that have been done (preferably with a timestamp) and when the device comes back online, transmits them all for the server to deal with each change in sequence⁷. This approach may lead to a lot more conflicts to be resolved — every request might generate a conflict, so more requests means more potential conflicts.

There is one more issue that arises in this approach: when creating a new instance, it is assigned an ID (locally) for unique identification. Later, but still in offline, the instance might be updated — this change will be recorded as a PUT with the generated ID. When the synchronization to server starts and the POST is transmitted, the server may assign it an ID based on its ID-generation strategy and this ID will probably be different than the locally generated one. The subsequent PUT would then have an invalid ID. To resolve the issue, the client needs to keep a relation of local and server IDs, much like a translation table, and use it to modify any subsequent calls that reference the instance.

Despite these limitations and added complexity, there are several reasons to prefer this approach:

1. Auditability: each change is recorded, so there is full control over who did what and when it has been done.
2. Granularity: when one change out of ten causes a conflict, in the operation-based approach it is clear which one it is, while the state-based approach will diffuse it.
3. Business logic: it may not be possible to implement the state-based approach in some cases due to functional requirements — when a change of an attribute has its meaning. Imagine a package delivery app where the package has a state (in warehouse, in delivery, delivered). The courier's phone is currently

7. Optimally, the server also accepts the client's timestamp and records the change to have been done in the time of the timestamp, not the request time. Therefore, the real timing of changes, creation, or deletion is also recorded.

not connected to the internet as he picks up the package from a warehouse and delivers it to the customer. It is only after that when their phone finally connects to the internet. In state-based approach, the package would go straight from *inWarehouse* state to *delivered* state — which is probably not the expected behaviour.

1.2.5 Offline authentication, authorization

Application usability strongly relies on the user remaining authenticated while offline — they cannot be automatically signed out when their token expires. Either the token-checking logic has to be suspended while offline, or the server response to such request must be mocked to a successful call.

But, remaining signed into the application while offline may cause a security issue — the local data can be sensitive and without authorization, an intruder that has taken over the device, even after a long time trying, can see the data unauthorized, provided that the device remains offline. This risk can be mitigated by, for example, signing the user out after a specific amount of time in offline, but that has to be agreed-upon from the business side.

Yet there is still an additional issue. Even when the user is signed out from the application, the local data is still present in the storage of the device. In case this poses a considerable security issue in the customer's context, the data must be encrypted by the application.

2 Technical approaches

This chapter explores different technical approaches to employ when designing an enterprise solution in the context outlined in the previous chapter.

2.1 Frameworks

First, there is the crucial decision over the base architecture of the application, that is, the language and framework to be used. This choice has a lot of implications for later design choices — for example, lack of certain features in certain frameworks may mean restricted choice of solutions for a particular challenge.

Since the topic at hand is mobile applications, naturally the first thing that comes to mind are mobile-specific (mobile-only) languages and frameworks. These can be split into native approaches such as Kotlin or Swift, and cross-platform ones such as React Native or Flutter.

But there has been a strong tendency to favor browser-based solutions over native ones in recent history, as Jon Henshaw[3] comments. He argues that Google has been one of the biggest driving forces on the market and that the sheer market share of their browser, Google Chrome, can in part be attributed to this transformation. The architecture in question is SPA (single page application), specifically in form of PWAs (progressive web applications).

Tools that Henshaw mentions, like Google Meet, YouTube or Canva, are only a fraction of the sheer quantity of applications that use this web-based architecture. Spotify, Netflix, or applications like BandLab¹ — a fully-featured in-browser music creation platform — after all, even Microsoft's office suite has been migrated to the browser with the Office 365 package and as Bill Doll[4] explains, Microsoft has been heavily investing to make their office suite web-first for quite a few years now, even prioritising the web variant over the native one for new feature releases.

1. <https://www.bandlab.com/>

2.1.1 Native approaches

The main argument to use native solutions is the closeness to the operating system. That means, the fewer layers the better performance, smoother UX (user experience) and better system integration. But it may also bring greater complexity and management.

Advantages

Performance is among the most prominent reasons to go native. Philipp Horn; Rituparna Bose; Christian Becker[5] performed a qualitative study that compared Native Applications with their web-based counterparts, finding that there is a statistically significant advantage of native apps in not only performance (startup time, responsiveness), but also power consumption, resource consumption (such as CPU² or memory) and network traffic. Therefore, when the application is performance-demanding, the native approach should not be overlooked.

There is a related advantage of native applications that enhances the user's perception of performance — the UI (user interface) / UX integration into the system. Since the applications are native, they respect the animations, navigation, and overall user experience of the system they are running on. They will be very well optimized for the device and may even be coupled with haptic feedback, system shortcuts or widgets.

The last notable strength of this architecture that should be discussed is operating-system-level integration. Native apps can communicate with the device's hardware and core features significantly better than a browser application ever will. That includes all the sensors that a mobile device has — camera, Bluetooth, NFC (Near Field Communication), Accelerometer/Gyroscope, ambient light, temperature sensors, and a lot more. All of these can be directly integrated into a native app. Their data are presented through respective system APIs in (more or less) raw form. There is little abstraction and maximum control.

2. Cetral Processing Unit

2. TECHNICAL APPROACHES

This area is not only limited to hardware itself, there is the OS (Operating System) layer as well — managing foreground and background jobs, having full control over OS notification channels, battery optimizations, or native sharing options and integration with other native applications. To conclude, when the application needs to have a deep hardware or operating level integration, the native approach plays its prime.

Disadvantages

On the other hand, there certainly are disadvantages to this approach. The greatest one being the fact that when multiple operating systems need to be supported, there have to be multiple codebases in use — each for one platform.

Typically, that would mean Kotlin for Android support and Swift for iOS support. There is quite a limited extent to which certain logic can be shared between different languages and frameworks. That means a lot more cost to develop and maintain two application versions instead of one. The applications also need to be published and then updated by respective app stores, so the deployment process can be a little more difficult.

Additionaly, there may also be a need to support Windows mobile devices as well, not even mentioning Ubuntu mobile and other alternatives.

The other disadvantage that should be talked about comes from the very same reason that is mentioned earlier as an advantage — the closeness to the OS and the level of control over it, or in other words, the absence of abstracting layers between the application and the OS. Less abstraction leads to heavier reliance on specific OS versions and flavors.

For example, every Android version on every phone brand differs slightly and to ensure the app compliance with a wide variety of supported devices, the testing needs to be extensive and the development might have to include a lot more specific case handling.

Kotlin

What is often referenced as Kotlin is in fact Kotlin + Android SDK (software development kit) + Jetpack Compose. Kotlin³ is a programming language developed by JetBrains. It was designed to be a modern yet mature programming language that takes a lot of inspiration from languages like Java, C# or Scala[6]. One of its strong points is that it is fully interoperable with Java, which makes migration as painless as possible.

Google states that Kotlin provides better expressiveness, conciseness and safety over Java, which are the main reasons of why in 2019, Google announced that Android development will be Kotlin-first[7]. According to the Stack Overflow Developer Survey from 2024[8], developers enjoy working with Kotlin more than with Java, which is shown by the admiration index where Kotlin surpasses Java by almost 15% (47.6% for Java, 60.9% for Kotlin).

Calling the whole stack just ‘Kotlin’ is the norm in the developer community not only because of it being the official language chosen by Google, but also because the main usage of the Kotlin language is Android development, as can be seen in the JetBrains 2021 Developer Ecosystem Report[9].

To comment on the other parts of the Android development stack, the Android SDK is a layer between the Android OS and the Kotlin code that enables developers to communicate with system resources in a form of an API (application programming interface) together with build tools or an emulator for testing. Jetpack Compose is a UI toolkit for Android applications. It is the recommended way to design and build the UI by Google[10].

Swift

Swift⁴ for iOS is what Kotlin is for Android. Developed by Apple, designed to replace Objective-C, much like Kotlin was designed to replace Java, it is the official language for its respective OS. Its main

3. <https://kotlinlang.org/>
4. <https://www.swift.org/>

focus is on speed, expressiveness and safety while providing interoperability with C and C++. According to the aforementioned Stack Overflow Developer Survey[8], Swift is even more admired amongst developers with 63.3% admiration index, while virtually obliterating Objective-C with its 26.3%.

In the technology stack, similarly to the Android ecosystem, Apple SDK provides the above-OS layer, and SwiftUI is the counterpart to Jetpack Compose.

Talking about comparisons between Kotlin and Swift, performance is the first that might come to mind. According to a set of benchmarks⁵, Kotlin seems to generally complete tasks faster when coupled with JVM (Java Virtual Machine). There are other technical differences such as Kotlin's JVM-powered Garbage Collection Approach (GCA) versus Swift's Automatic References Counting (ARC) for memory management[11] and other caveats, yet there is not much sense in such comparisons as the languages are designed for completely different environments.

There is one direct comparison, however, that is to be made - the syntax of the languages, revealing their expressiveness. Two code snippets of a simple class follow, the former in Swift and the latter in Kotlin. As expressiveness is subjective, it is up to the reader to decide the superior one.

Listing 2.1: Swift example

```
class Pet {  
    var name: String  
    var age: Int  
    var sound: String?  
  
    init(name: String, age: Int, sound: String?  
         = nil) {  
        self.name = name  
        self.age = age  
        self.sound = sound  
    }  
}
```

5. <https://programming-language-benchmarks.vercel.app/swift-vs-kotlin>

```

func getSoundMessage() -> String {
    return (sound != nil) ? "\$(sound!)" : "
        I do not make any sound..."
}

func printGreeting() {
    let soundMessage = getSoundMessage()
    print("Hello, my name is \$name and I
        am \$age years old. \$soundMessage")
}
}
  
```

Listing 2.2: Kotlin example

```

class Pet(
    var name: String,
    var age: Int,
    var sound: String? = null
) {

    fun getSoundMessage(): String = sound?.plus(
        " !") ?: "I do not make any sound..."

    fun printGreeting() {
        val soundMessage = getSoundMessage()
        println("Hello, my name is \$name and I
            am \$age years old. \$soundMessage")
    }
}
  
```

2.1.2 Cross-platform approaches

As has been highlighted in the last section, one of the main drawbacks of the native approaches is that when multiple operating systems must be supported, the same functionality has to be implemented in multiple codebases. That is a major reason not to elect them as the ideal

solution — and that is where multiplatform approaches come into play. There are two main options on the market to choose — Flutter or React Native.

A layer of isolation between the application and the OS is introduced that abstracts the OS-specific differences such as system APIs or other SDK functionality — which carries with it its pros and cons.

Advantages

The most significant advantage is having a single codebase, which reduces costs dramatically. The layer of abstraction over the OS ensures that no matter the specificity of the OS, there will be common functionality that is expected to work on all of them. This layer communicates directly with the OS itself, so hardware availability, like access to camera, filesystem, or sensors, is still supported.

Disadvantages

Despite this approach solving the biggest problems with native solutions, it is far from perfect. According to research done by Ihor Demedyuk; Nazar Tsybulskyi[12], Flutter and especially React Native have poorer performance when compared to their native counterparts.

The abstraction over more operating systems also means the control over specific functionality may not be so precise. For example, there might be some camera controls or sensors available only on one OS that may get dropped during the abstraction or must be handled with platform-specific code. New features may also come with a delay as they have to be layered over first instead of just implementing the system API in Kotlin or Swift.

Also, the app still needs to be distributed over multiple system app stores, which may cause delay in delivering new versions of the software.

Flutter

Flutter⁶ is a framework developed by Google that can be compiled to native code and be used on Android, iOS, desktop or web[13]. It uses its own rendering engine developed with Skia (a 2D rendering engine) and Impeller (iOS rendering engine), which means that the UI will be uniform on all platforms regardless of the platform-specific UI or animations. While that may improve brand consistency, it may suffer from OS inconsistencies.

Flutter uses Dart⁷ as its programming language. Developed by Google, Dart is not a broadly used language — in the 2024 Stack Overflow Developer Survey[8], it reached only 6% in popularity and 6.2% in desirability. Its admirability index is not anything special either with the score of 55%. It is a strongly typed, object oriented language that can be compiled to machine code, JavaScript, or WebAssembly.

In the aforementioned performance benchmark, Flutter achieves significantly higher performance in intensive tests. If performance is an important factor, but native approach is not to be considered, Flutter is the better option.

Flutter uses a tree-like widget structure for its UI — a widget is a single building block for any UI element. It resembles native-like syntax, meaning Kotlin or Swift developers should find it approachable. A short code snippet follows that illustrates the widget declaration — a simple clickable User Card that displays the user's name and profile picture; clicking it navigates the user to the user screen.

Listing 2.3: Flutter example — User Card

```
class UserCard extends StatelessWidget {
    final User user;

    const UserCard({Key? key, required this.user})
        : super(key: key);

    @override
```

6. <https://flutter.dev/>
7. <https://dart.dev/>

```
Widget build(BuildContext context) {
    return GestureDetector(
        onTap: () => Navigator.pushNamed(context,
            '/user'),
        child: Card(
            child: ListTile(
                title: Text(user.name),
                trailing: CircleAvatar(
                    backgroundImage: NetworkImage(user.
                        profilePictureUrl),
                    radius: 16,
                ),
            ),
        ),
    );
}
```

React Native

React Native⁸ is a cross-platform framework developed by Meta that can be used to develop mobile applications for both iOS and Android. It uses JavaScript (JS) or TypeScript (TS) as its programming language. This is one of the biggest advantages over Flutter as those languages are very popular, meaning there is a huge developer base and an abundance of tutorials, guides and forum pages to help the developer when encountering an issue.

Specifically, according to the Stack Overflow 2024 Developer Survey[8], Javascript is the most popular language with 62.3% popularity while TypeScript is fifth overall with 38.5% populartiy. They are also desired (39.8% — second place overall — for JS and 33.8% — fifth place overall — for TS) and TypeScript is also admired with 69.5% of admirability (JavaScript falls short at 58.3%).

8. <https://reactnative.dev/>

The JavaScript/TypeScript popularity also means there is a vast universe of third party libraries developed for them, meaning that when in need of a specific library, e.g. a barcode scanner, there is a good chance someone has built that library and it can be used in the project as a plug-and-play component.

React Native does not have a custom rendering engine like Flutter does, its code is rendered as native UI elements. This connection between the JavaScript and the native world is possible thanks to a “bridge”[14]. This bridge can also be used to write platform-specific native code, for example in Kotlin. This unique feature allows React Native applications to overcome the abstraction layer when needed, which may be very powerful.

React Native uses *components* as building blocks, much like Flutter uses widgets. They, much like in React, are JSX/TSX functions, therefore written in markup style. To illustrate the syntax difference compared to Flutter, a JSX component of a User Card follows:

Listing 2.4: React Native example — User Card

```
export default function UserCard({ user }) {
  const navigation = useNavigation();

  return (
    <TouchableOpacity onPress={() => navigation.
      navigate("User")}>
      <Card>
        <Card.Title
          title={user.name}
          right={() => (
            <Image
              source={{ uri: user.
                profilePictureUrl }}
              style={{ borderRadius: 16 }}>
            />
          )}
        />
      </Card>
    </TouchableOpacity>
```

```
) ;  
}
```

2.1.3 React PWA

React.js⁹ is a powerful JavaScript frontend library/framework used for building modern user interfaces. Developed by Facebook/Meta and published in 2013[15], it is by far the most popular frontend framework at 39.5%, followed by Next.js¹⁰ at 17.9%, which is a web framework built on top of React. It is also the most desired at 33.4% by a considerable margin and is reasonably admired at 62.2%, according to the Stack Overflow Developer Survey 2024[8].

From these numbers, it is safe to say that React is the single most used and important frontend framework on the current market. That means there is a massive amount of packages to be used for development in React and a multitude of guides, forums and other content about the framework which serves as an invaluable source of assistance during development. There is also a remarkable number of React developers on the market.

React uses JavaScript or TypeScript as the programming language (pros and cons of this choice are mentioned in the previous section about React Native), conceptually splitting the UI into singular components written in JSX (or TSX for TypeScript projects) syntax that serves as a JavaScript extension to write markup code — similar to HTML (Hypertext Markup Language), but enriched with JS elements — inside of JS functions[16]. For styling purposes, pure CSS (Cascading Style Sheets¹¹) or any of its frameworks like Tailwind¹² may be used.

An example of the (by now well-known) User Card follows:

Listing 2.5: React + Tailwind example — User Card

-
- 9. <https://react.dev/>
 - 10. <https://nextjs.org/>
 - 11. <https://developer.mozilla.org/en-US/docs/Web/CSS>
 - 12. <https://tailwindcss.com/>

```
export default function UserCard({ user }: Props) {
  const navigate = useNavigate();
  return (
    <div
      className="flex flex-row justify-between rounded-xl shadow-lg p-4"
      onClick={() => navigate("/user")}
    >
      <span className="text-lg font-semibold">{user.name}</span>
      <img src={user.profilePictureUrl} className="rounded-full" />
    </div>
  );
}
```

PWA

A PWA (Progressive Web Application)¹³ is a web application built as a capable standalone installable application enhanced with modern browser system APIs and offline capabilities[17].

A PWA can be installed straight from the browser through a context menu, meaning there is no need to handle app store publication or different builds for different platforms. There is only a singular codebase and a singular deployment — the “browser one” — that also provides the option to install the application on any platform, transforming it to a natively-feeling one.

If needed, there is also a packaging option called PWABuilder¹⁴ that can produce a application package that can be distributed to several app stores like Google Play Store.

An installed PWA can be pinned to a mobile home screen, PC desktop, or a taskbar. It can be managed inside the OS, set up to use

13. <https://web.dev/explore/progressive-web-apps>

14. <https://pwabuilder.com/>

push notifications, registered to accept content from other system applications, set up to override browser keyboard shortcuts, or chosen to be the default application for opening files of certain types.

A PWA also makes web applications offline-ready. It has become a user-expected standard in the last couple of years for applications to provide some form of content even when offline. In the traditional Web Application context, that is not possible — the web application frontend itself is provided to the browser from the server, and when the request to load a page fails (because of network issues or server unavailability), the browser just displays an error page. A PWA enables caching strategies handled by service workers.

A Service Worker is a middleware residing in the browser that acts as a network proxy between the client and the server[18]. The service worker has a developer-specified scope of network requests that it handles. Depending on the strategy chosen for each section of endpoints, it can handle them accordingly — serving fresh or cached, serving a custom-made data, or, for example, serving from cache and fetching a fresh version of data in the background and updating the cache.

Advantages and Disadvantages

When weighing pros and cons of progressive web applications, one can view it is the polar opposite of the native approach — while the native solution is as close to the hardware and OS, this approach is as abstracted as possible, layering over the OS with not only a cross-platform layer over the system APIs like Flutter does, but a whole browser layer over it as well.

There is a certain level of advantages coming from this fact — with the solution practically detached from the OS, the developer only needs to care about the browser capabilities. Features like camera support include the enumeration of front or back facing cameras, abstracting from PC webcams to multiple lenses on smartphones. Another example may be a file picker, handled by the browser, bringing up the correct dialog or window on any system and handling the file upload itself.

The browser abstraction layer also means that the PWA can only do things that the underlying browser can — opposite to native, the browser cannot influence the battery management, reach sensors like fingerprint sensor, WiFi details or any health-related sensor data.

The extent to which other system data or functionality is available is also restricted. Bringing up the camera once again, the in-browser camera capabilities are limited by the WebRTC¹⁵ protocol stack — the camera quality is typically heavily reduced and its raw output is not processed by the OS like it would be in the native camera application on the device. That may lead to user frustration — why is the quality so much worse compared to taking the photo outside of the application when the hardware is the same?

All the abstraction over the OS also means that the performance will be poorer when compared to native solutions when doing resource-intensive tasks. Yet the PWA may try to tackle this using smart caching strategies.

2.1.4 Conclusion

There are a few clear takeaways from the selected approaches to be taken when choosing the main application framework:

1. When performance in resource-intensive tasks is crucial, there is no better solution than native. Yet for “normal” usage, which is the expected one when talking about the context of enterprise mobile applications, the user may not feel the speed difference.
2. When the application’s focus is on utilizing the device’s hardware capabilities like NFC or heart-rate sensor, the correct choice is also to go native.
3. If possible and in accordance with the functional requirements, when needing to support more operation systems, going cross-platform may save a lot of resources thanks to only maintaining a single codebase.

15. <https://webrtc.org/>

4. When there is an option to go the PWA route, it should always be considered as it may be the cheapest and smoothest one thanks to a large developer base and a plethora of libraries to be used to enhance the application.
5. PWA is the easiest to distribute, fastest to install and most consistent over multiple platforms. But being the most abstracted one, precise access to various system resources is restricted.

2.2 Local data management

As chapter one explains, keeping track of local data and its changes in offline mode is a crucial issue to tackle. There are two magnitudes to be chosen for resolving this problem — faking just the server or also faking a database.

2.2.1 Request intercepting and processing

The first option is to build a custom proxy-like layer serving as a form of middleware between the client and the server. This layer intercepts all the mutating requests. The data getting (GET) requests are handled separately, for example in the form of cache. The POST, PUT, PATCH, and DELETE requests are stored as pending operations and a success state (200, 201 etc.) is returned to the client as if the request was successfully processed by the server. Later all the pending operations are synchronized one by one, in order, with the server.

Client validation

This *optimistic* UI update strategy also means that the data validation must always be done on the client and data consistency must be enforced there — there is no server to handle it, the client needs to ensure that the request have the highest possible probability of succeeding once they are actually transmitted through the network and processed by the server.

Importance of this principle grows with any inter-endpoint or stateful dependencies. As an example, let's imagine an entity with several states and a lot of attributes. In each state, a certain subset of attributes is displayed to be filled. A change of state of an instance requires the instance (or its related/subordinate entities) to be filled with a valid values for the currently displayed set of attributes. The state change opens another subset of data to be filled, depending on the values of the previously submitted values.

That is a *dependency path* — a sequence of unique steps dependent on each other. If the user fills the first step incorrectly and the client fails to enforce the validity to the same extent that the server eventually will, each subsequent update to this instance must be thrown away during synchronization and the user, because the validity of the first step is compromised, effectively compromising all the following steps — the state change was not valid and therefore all subsequent dependent actions will not be either.

Incorporating mutations into local data

Assuming the validation is handled correctly, there is still one aspect that needs to be talked about — displaying the mutated data correctly. Earlier in this section, it was stated that the cached version of data is used for presentation. That fails as soon as there is a need to display mutated (updated) data in the UI — the pending operations need to be incorporated into the data.

The solution is to proxy also the GET requests and include any pending changes into the data presented by the client. As each entity is identified by its name (and therefore a specific endpoint subtree) and every instance is identified by its ID, it effectively means fetching the instance data from the cache and searching through the pending mutations to filter those that affect this instance — in other words, search for match in endpoint name and instance ID. All found mutations are then applied to the fetched (cached) data inside the middleware and presented to the client.

This intercepting approach fails to scale well for request-frequent or data-intensive applications as each GET request must be heavily

processed and a structure of pending requests searched through repetitively. The upside is that the client does not have to know about being offline at all as the middleware acts as an optimistic server — the client then fetches and mutates data as if the device was online and all is handled by the intercepting layer.

POST - PUT chain

In the previous section, an assumption about each instance being uniquely identified by an ID was made. There is one edge case that invalidates this assumption — creating an instance while offline. Traditionally, the server assigns IDs to new instances in a way that respects the integrity constraints of the database. The client, however, cannot reliably assign IDs on its own and synchronizing the POST request afterwards will lead to the ID being overwritten by the server, rendering all potential subsequent PUT/PATCH/DELETE requests to the same entity invalid as the ID would not match.

The simplest option is to prohibit the POST-PUT chain completely. If it makes sense in the context of a particular application where the vast majority of requests are mutating server data, the few use-cases that are creating and subsequently updating an instance may be transformed to simply remove the initial POST from the pending operations and creating an entirely new one. This approach is the simplest one, but when creation/update timestamping is crucial, it is not a viable one.

The second option is to generate the ID on POST on the client and later during the synchronization map the client-generated ID to the real server-assigned one and then alter all subsequent calls to this instance to respect the correct ID.

2.2.2 Local database

The second, more robust, scalable, but also complex and demanding solution is to manage a local database on the client. The support for local databases is present in all frameworks mentioned in this thesis —

2. TECHNICAL APPROACHES

from the Room database in Kotlin¹⁶ to the browser's IndexedDB for PWAs¹⁷.

The local database schema needs to closely mirror the server one and must implement the same integrity constraints, keys, or unique constraints to mitigate synchronization conflicts. The relevant data is fetched and stored to the local database while online and then, when offline, the application only works with the local database to access and mutate data, no requests are attempted. Relevant instances may be retrieved, created, updated or deleted from the database.

During the synchronization, the local tables must be merged with the server ones, writing through the data changes that occurred while offline. The changed/created/deleted instances from the offline session would be flagged accordingly and requests created that would reproduce said data changes.

The obvious downside to this approach is its complexity and the need to adhere closely to the server architecture. Any changes on the backend must be properly mirrored on the frontend and the thickness of the client therefore grows, meaning more expensive development and maintenance. On the other hand, the robustness of such solution is incomparable to the previous one, because there is strong consistency in the data (all relevant data is always updated properly, which may not be true in the first approach if there are complex dependencies across entities).

2.2.3 Hybrid solutions

There is also a “hybrid solution” — one that uses the database, that is controlled by the middleware. Effectively, the application is completely unbothered by being offline, but the requests sent to the server are intercepted by the middleware, processed, translated to local database queries, and also stored to a pending operations queue. The GET requests are transformed to retrieve data from the local database, because all changes are reflected in it.

16. <https://developer.android.com/training/data-storage/room>

17. https://developer.mozilla.org/en-US/docs/Web/API/IndexedDB_API

Upon getting back online, the synchronization is as simple as in the first approach — sending the requests to the server in the order of interception. Combining the strengths of both approaches, the cost is in the increased complexity.

Now, effectively, the backend together with the database is completely mocked in the client. This significantly increases costs and maintainability of such a solution and therefore should be used in specific scenarios only, especially when there are complex data dependencies and high independence requirements (the client must remain offline for longer periods of time).

2.3 Synchronization

Regardless of the choice of the local data handling and therefore synchronization method, there may always arise conflicts. Conflicts arise when the offline and online data are both changed and therefore one change may overwrite the other.

To illustrate the approaches described in this section on a practical scenario, let's assume this situation: there are two users, Alice and Bob. They both work in the same time with the same system and they interact with the same data, or at least a portion of the data they work with is shared. Alice works online, while Bob is in offline mode. After some time, Bob reconnects to the internet and tries to synchronize his changes, creating conflicts on the shared data.

2.3.1 Mitigation

The simplest yet most restrictive approach is for Bob to lock preemptively the instances that he intends to interact with while offline, effectively mitigating the risk of any conflict.

An instance-granular locking mechanism must be implemented, and Bob, before going offline, has to be able to mark the instances that he needs in offline and lock them. Optionally, instances might be locked together with all their “subinstances” — that is, instances that have a $n : 1$ relation with said instance.

There are, unfortunately, several problems to this approach — firstly, Alice is blocked from working with the instances completely. Even though they might not overwrite each other's attributes as the attribute subsets do not overlap, she shall not interact with the instances at all. Secondly, Bob has to know what data he will interact with in advance.

And thirdly, after he is done working with the data, he must always unlock the instances on the server — failure to do so will block any other user from interacting with it over potentially a prolonged period of time. Assuming Bob went on holiday for a week and did not bring his work device with him, the only option for his colleagues to work with the data is to have a system administrator that can forcefully unlock the instances, effectively scrapping Bob's unsaved changes.

2.3.2 Duplication

2.3.3 Attribute timestamping

2.3.4 CRDT

2.3.5 Advanced synchronization management

(approaches based on advanced methods like GIT etc)

3 Implementation

3.1 Functional and non-functional requirements

3.2 Architecture, technology stack

3.3 Service Worker

popsat ktery strategie realne pouzivam

popsat ruzne strategie a jejich vyuuziti/vyhody nevyhody?

3.4 Axios Interceptor

3.5 Local data management

3.6 Synchronization

zminit ze musim posilat nejaky timestamp kdy se to stalo v offlinu
realne, nejen kdy se odeslal ten request na server pri synchronizaci, to
neni relevantni.

4 Testing, issues, and future work

Conclusion

Bibliography

1. MORETTI, Francisco. Next.js Revolution — The framework popularity rises in 2023 [online]. 2023 [visited on 2025-09-27]. Available from: <https://medium.com/@franciscomoretti/next-js-revolution-the-framework-popularity-rises-in-2023-a8ebc52d673a>.
2. ZHAO, Jing; ZHANG, Fan; ZHAO, Chao; WU, Gang; WANG, Haitao; CAO, Xinyu. The Properties and Application of Poisson Distribution. 2020, vol. 1550, no. 3, p. 032109. Available from doi: 10.1088/1742-6596/1550/3/032109.
3. HENSHAW, Jon. *Why Web Apps Are Quietly Replacing Native Apps on Phones and Computers*. Coywolf News, 2024. Available also from: <https://coywolf.com/news/productivity/why-web-apps-are-quietly-replacing-native-apps-on-phones-and-computers/>. Accessed: 2025-10-20.
4. DOLL, Bill. *Why You Should Use Office for the Web*. Microsoft 365 Blog, 2019. Available also from: https://techcommunity.microsoft.com/blog/microsoft_365blog/why-you-should-use-office-for-the-web/567060. Accessed: 2025-10-20.
5. HORN, Philipp; BOSE, Rituparna; BECKER, Christian. Native vs Web Apps: Comparing the Energy Consumption and Performance of Android Apps and their Web Counterparts. *arXiv preprint arXiv:2308.16734*. 2023. Available also from: <https://arxiv.org/abs/2308.16734>. Accessed: 2025-10-20.
6. GOOGLE LLC – ANDROID DEVELOPERS. *Kotlin and Android* [<https://developer.android.com/kotlin/overview>]. 2024. Accessed: 2025-10-23.
7. GOOGLE LLC – ANDROID DEVELOPERS. *Android's Kotlin-first approach* [<https://developer.android.com/kotlin/first>]. 2024. Last updated June 27, 2024. Accessed: 2025-10-23.
8. STACK OVERFLOW. *Stack Overflow Developer Survey 2024*. 2024. Available also from: <https://survey.stackoverflow.co/2024/>. Accessed: 2025-10-23.

BIBLIOGRAPHY

9. JETBRAINS. *The State of Developer Ecosystem 2024* [<https://www.jetbrains.com/lp/devcosystem-2024/>]. 2024. Accessed: 2025-10-23.
10. GOOGLE LLC – ANDROID DEVELOPERS. *Jetpack Compose — Android's recommended modern UI toolkit* [<https://developer.android.com/compose>]. 2025. Accessed: 2025-10-23.
11. JAFTON. *Kotlin vs. Swift: How Are They Different?* Jafton, 2025. Available also from: <https://www.jafton.com/insights/kotlin-vs-swift>. Accessed: 2025-10-27.
12. DEMEDYUK, Ihor; TSYBULSKYI, Nazar. *Flutter vs Native vs React-Native: Examining Performance* [online]. inVerita, 2020. [visited on 2025-10-20]. Available from: <https://medium.com/swlh/flutter-vs-native-vs-react-native-examining-performance-31338f081980>.
13. FLUTTER TEAM. *Flutter FAQ* [online]. 2025. [visited on 2025-10-27]. Available from: <https://docs.flutter.dev/resources/faq>.
14. BIRD, Shane. *React Native: All You Need To Know* [online]. Blott Studio, 2022. [visited on 2025-10-27]. Available from: <https://www.blott.com/blog/post/react-native-all-you-need-to-know>.
15. REACT TEAM. *Versions – React* [online]. Meta Platforms, Inc., 2025. [visited on 2025-10-27]. Available from: <https://react.dev/versions>.
16. REACT TEAM. *Writing Markup with JSX* [online]. Meta Platforms, Inc., 2025. [visited on 2025-10-27]. Available from: <https://react.dev/learn/writing-markup-with-jsx>.
17. LEPAGE, Pete; RICHARD, Sam. *What are Progressive Web Apps?* [online]. Google LLC, 2020. [visited on 2025-10-27]. Available from: <https://web.dev/articles/what-are-pwas>.
18. FIRTMAN, Maximiliano. *Service workers* [online]. Google LLC, 2021. [visited on 2025-10-27]. Available from: <https://web.dev/learn/pwa/service-workers>.

A Electronic attachments